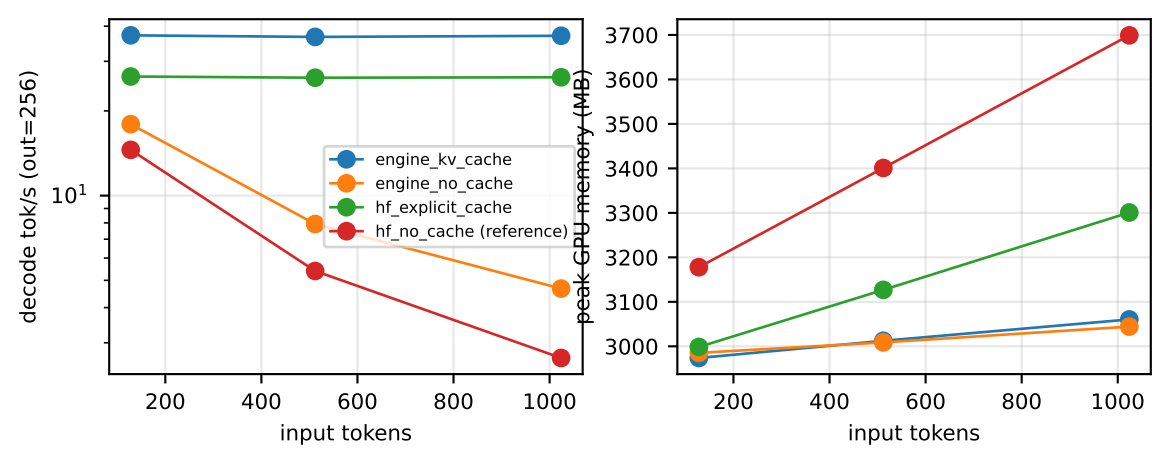
decode tok/s (out=256)
10^1
input tokens
200
400
600
800
1000
engine_kv_cache
engine_no_cache
hf_explicit_cache
hf_no_cache (reference)
peak GPU memory (MB)
3700
3600
3500
3400
3300
3200
3100
3000
input tokens
200
400
600
800
1000